

```
In [32]: # Import required packages
import numpy as np
import time
import cv2
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression

import tensorflow as tf
import numpy as np
import pandas as pd
from tensorflow.keras.callbacks import ModelCheckpoint, ReduceLROnPlateau
from tensorflow.keras.layers import Input, Dense, BatchNormalization, Flatten, MaxPooling2D, Activation, GlobalMaxPool2D, GlobalAvgPool2D, Concatenate, Multiply, Dropout, Subtract
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
from tensorflow.keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img
from tensorflow.keras.optimizers import SGD, Adam, RMSprop, Nadam
from sklearn.utils import shuffle
from sklearn.model_selection import train_test_split
```

```
In [33]: print(f"This notebook uses TensorFlow Version {tf.__version__}")
print("And Python Version:")
!python --version
```

This notebook uses TensorFlow Version 2.6.0
And Python Version:
Python 3.6.8 :: Anaconda, Inc.

Results

Model I Results:

Noisy Train/Test Set Accuracy: 24.42%, 22.32% Model I Clean Labels Loss, Accuracy: [1.744310975074768, 0.5378000140190125]

Clean Image Accuracy: 53.78%

Training Time: 1013.9775972366333 Seconds

Model II results:

Train/Test set Accuracy: 56.99%, 56.26%

313/313 [=====] - 9s 29ms/step - loss: 1.2218 - accuracy: 0.5703 Model II Clean Labels Loss, Accuracy: [1.2217928171157837, 0.5702999830245972]

Clean Image Accuracy: 57.03%

Train Time: 1468.9489738941193 + 587.4192531108856 = 2056.37

1. Load the datasets

For the project, we provide a training set with 50000 images in the directory `../data/images/` with:

- noisy labels for all images provided in `../data/noisy_label.csv`;
- clean labels for the first 10000 images provided in `../data/clean_labels.csv`.

```
In [34]: # [DO NOT MODIFY THIS CELL]

# load the images
n_img = 50000
n_noisy = 40000
n_clean_noisy = n_img - n_noisy
imgs = np.empty((n_img, 32, 32, 3))
for i in range(n_img):
    img_fn = f'../data/images/{i+1:05d}.png'
    imgs[i, :, :, :] = cv2.cvtColor(cv2.imread(img_fn), cv2.COLOR_BGR2RGB)

# load the labels
clean_labels = np.genfromtxt('../data/clean_labels.csv', delimiter=',',
                             dtype="int8")
noisy_labels = np.genfromtxt('../data/noisy_labels.csv', delimiter=',',
                              dtype="int8")
```

For illustration, we present a small subset (of size 8) of the images with their clean and noisy labels in `clean_noisy_trainset`. You are encouraged to explore more characteristics of the label noises on the whole dataset.

```

In [127]: # [DO NOT MODIFY THIS CELL]

fig = plt.figure()

ax1 = fig.add_subplot(2,4,1)
ax1.imshow(imgs[0]/255)
ax2 = fig.add_subplot(2,4,2)
ax2.imshow(imgs[1]/255)
ax3 = fig.add_subplot(2,4,3)
ax3.imshow(imgs[2]/255)
ax4 = fig.add_subplot(2,4,4)
ax4.imshow(imgs[3]/255)
ax1 = fig.add_subplot(2,4,5)
ax1.imshow(imgs[4]/255)
ax2 = fig.add_subplot(2,4,6)
ax2.imshow(imgs[5]/255)
ax3 = fig.add_subplot(2,4,7)
ax3.imshow(imgs[6]/255)
ax4 = fig.add_subplot(2,4,8)
ax4.imshow(imgs[7]/255)

# The class-label correspondence
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# print clean labels
print('Clean labels:')
print(' '.join('%5s' % classes[clean_labels[j]] for j in range(8)))
# print noisy labels
print('Noisy labels:')
print(' '.join('%5s' % classes[noisy_labels[j]] for j in range(8)))

```

Clean labels:

frog truck truck deer car car bird horse

Noisy labels:

cat dog truck frog dog ship bird deer



2. The predictive model

We consider a baseline model directly on the noisy dataset without any label corrections. RGB histogram features are extracted to fit a logistic regression model.

2.1. Baseline Model

```
In [35]: # [DO NOT MODIFY THIS CELL]
# RGB histogram dataset construction
no_bins = 6
bins = np.linspace(0,255,no_bins) # the range of the rgb histogram
target_vec = np.empty(n_img)
feature_mtx = np.empty((n_img,3*(len(bins)-1)))
i = 0
for i in range(n_img):
    # The target vector consists of noisy labels
    target_vec[i] = noisy_labels[i]

    # Use the numbers of pixels in each bin for all three channels as the features
    feature1 = np.histogram(imgs[i][:,:,0],bins=bins)[0]
    feature2 = np.histogram(imgs[i][:,:,1],bins=bins)[0]
    feature3 = np.histogram(imgs[i][:,:,2],bins=bins)[0]

    # Concatenate three features
    feature_mtx[i,:] = np.concatenate((feature1, feature2, feature3), axis=None)
    i += 1
```

```
In [36]: # [DO NOT MODIFY THIS CELL]
# Train a logistic regression model
clf = LogisticRegression(random_state=0).fit(feature_mtx, target_vec)

/Users/kerry.cook@ibm.com/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
/Users/kerry.cook@ibm.com/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.
  "this warning.", FutureWarning)
```

For the convenience of evaluation, we write the following function `predictive_model` that does the label prediction. **For your predictive model, feel free to modify the function, but make sure the function takes an RGB image of numpy.array format with dimension $32 \times 32 \times 3$ as input, and returns one single label as output.**

```
In [37]: # [DO NOT MODIFY THIS CELL]
def baseline_model(image):
    '''
    This is the baseline predictive model that takes in the image and re
    turns a label prediction
    '''
    feature1 = np.histogram(image[:, :, 0], bins=bins)[0]
    feature2 = np.histogram(image[:, :, 1], bins=bins)[0]
    feature3 = np.histogram(image[:, :, 2], bins=bins)[0]
    feature = np.concatenate((feature1, feature2, feature3), axis=None).
    reshape(1, -1)
    return clf.predict(feature)
```

2.2. Model I

For model I, we use a basic CNN structure: two 2D convolutional layers, a max pooling layer, a flatten layer, a dense layer and the classification layer.

For the optimizer, we use Nadam and the learning rate is 0.001.

We use data augmentation in order to reduce overfitting. The data augmentation also increases the amount of data as it adds modified copies of the original data.

Split Data

```
In [38]: # shuffle the images and split the data into training and validation set
(0.9 and 0.1)
shuff_imgs, target_vec = shuffle(imgs, noisy_labels, random_state=0)
img_train, img_test, y_train, y_test = train_test_split(shuff_imgs, targ
et_vec, test_size=0.10, random_state=42)
# one-hot encoded vectors
y_train = np.eye(10)[y_train]
y_test = np.eye(10)[y_test]
```

```

In [39]: def modelI():

    # Simple CNN with 2 convolutional layers, max pooling, flatten layer
    and dense layer
    model = Sequential()

    model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32,
3)))
    model.add(Conv2D(32, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Flatten())
    model.add(Dense(256, activation='relu'))
    model.add(Dense(10, activation='softmax'))

    # COMPILE
    opt= Nadam(
        learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, name=
"Nadam")

    # compile
    model.compile(optimizer=opt,
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])

    return model

def train_model(model, img_train, y_train, img_test, y_test, output_fn,
epochs = 10 ):

    # generate image data with data augmentation
    train_gen = ImageDataGenerator(
        featurewise_center=True, # set the mean of the inputs to 0 over
the dataset
        featurewise_std_normalization=True, # divide the inputs by stand
ard deviation of the dataset
        rotation_range=20, # degree range for random rotations
        width_shift_range=0.2, # the fraction of total width
        height_shift_range=0.2, # the fraction of total height
        horizontal_flip=True) # flip the inputs horizontally randomly

    train_gen.fit(img_train)

    test_gen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True
    )
    test_gen.fit(img_train)

    # save the weights
    file_path = f"../output/{output_fn}"
    checkpoint = ModelCheckpoint(file_path, monitor='val_accuracy', verb
ose=1, save_best_only=True, save_weights_only=True, mode='max')
    callbacks_list = [checkpoint]

    # fits the model on batches with data augmentation:

```

```
model.fit(train_gen.flow(img_train, y_train, batch_size=128),
          validation_data=train_gen.flow(img_test, y_test, batch_size
=12),
          callbacks=callbacks_list,
          epochs=epochs)

return model, test_gen

def model_I(image):
    '''
    This function should takes in the image of dimension 32*32*3 as input
    and returns a label prediction
    '''
    # load the model weights
    model1 = modelI()
    model1.load_weights("../output/modelI.h5")

    # predict
    pred = model1.predict(data_genI.flow(image))
    pred_class = np.argmax(pred)

    return pred_class
```

Train the CNN for 6 epochs - would train for more, but only saw 2-3% gains, so reduced to 6 for the runtime reduction.

```
In [40]: # record the computational time
start = time.time()
# train model 1
model = modelI()
model, data_genI = train_model(model, img_train, y_train, img_test, y_test, "modelI.h5", epochs = 6)
end = time.time()
```

Epoch 1/6

352/352 [=====] - 221s 622ms/step - loss: 2.2752 - accuracy: 0.1603 - val_loss: 2.2522 - val_accuracy: 0.1888

Epoch 00001: val_accuracy improved from -inf to 0.18880, saving model to ../output/modelI.h5

Epoch 2/6

352/352 [=====] - 143s 404ms/step - loss: 2.2426 - accuracy: 0.1904 - val_loss: 2.2417 - val_accuracy: 0.1930

Epoch 00002: val_accuracy improved from 0.18880 to 0.19300, saving model to ../output/modelI.h5

Epoch 3/6

352/352 [=====] - 134s 381ms/step - loss: 2.2300 - accuracy: 0.2029 - val_loss: 2.2325 - val_accuracy: 0.2104

Epoch 00003: val_accuracy improved from 0.19300 to 0.21040, saving model to ../output/modelI.h5

Epoch 4/6

352/352 [=====] - 132s 374ms/step - loss: 2.2213 - accuracy: 0.2104 - val_loss: 2.2327 - val_accuracy: 0.2024

Epoch 00004: val_accuracy did not improve from 0.21040

Epoch 5/6

352/352 [=====] - 186s 527ms/step - loss: 2.2153 - accuracy: 0.2165 - val_loss: 2.2309 - val_accuracy: 0.2062

Epoch 00005: val_accuracy did not improve from 0.21040

Epoch 6/6

352/352 [=====] - 157s 446ms/step - loss: 2.2111 - accuracy: 0.2214 - val_loss: 2.2146 - val_accuracy: 0.2254

Epoch 00006: val_accuracy improved from 0.21040 to 0.22540, saving model to ../output/modelI.h5

```
In [41]: print( f"Total Model I training time: {end-start}")
```

Total Model I training time: 1013.9775972366333


```
In [43]: train_metrics = model.evaluate(data_genI.flow(img_train, y_train))
test_metrics = model.evaluate(data_genI.flow(img_test, y_test))

print(f"Model I Training Loss, Accuracy: {train_metrics}")
print(f"Model I Testing Loss, Accuracy: {test_metrics}")

1407/1407 [=====] - 44s 31ms/step - loss: 2.18
63 - accuracy: 0.2442
157/157 [=====] - 3s 17ms/step - loss: 2.2104
- accuracy: 0.2232
Model I Training Loss, Accuracy: [2.1862809658050537, 0.244222223758697
5]
Model I Testing Loss, Accuracy: [2.2104125022888184, 0.223199993371963
5]
```

```
In [42]: # clean images and labels
img_cl = imgs[:10000]
y_cl = np.eye(10)[clean_labels]

# estimate the model accuracy using clean labels
cl_metrics = model.evaluate(data_genI.flow(img_cl, y_cl))

print(f"Model I Clean Labels Loss, Accuracy: {cl_metrics}")

313/313 [=====] - 6s 20ms/step - loss: 1.7443
- accuracy: 0.5378
Model I Clean Labels Loss, Accuracy: [1.744310975074768, 0.537800014019
0125]
```

2.3. Model II

For Model II, we train a label cleaning network that follows a similar architecture as the paper. We used a pre-trained CNN (VGG16) for the base network, and tried to match the rest of the architecture to the paper. We make the last layer of VGG16 to be trainable in order to avoid overfitting.

The label network is only trained for 6 epochs, as it is time intensive, but performance could be increased by training for more epochs.

We then use the label cleaning network to predict new labels for the 40000 noisy images, and use the new labels along with the 10000 clean labels to retrain a new CNN that has the same architecture as Model 1. Overall accuracy increased.

Label Cleaning Network

```
In [44]: #Get both clean and noisy labels for the first 10,000 images
clean = np.eye(10)[clean_labels]
noisy = np.eye(10)[noisy_labels[:10000]]
clean_imgs = imgs[:10000]/255
```

```

In [45]: import tensorflow.keras.backend as K
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Lambda

#Custom loss function for comparing predicted class to clean label used
for training label network
def label_loss(y_true, y_pred):
    # L1 distance between true labels and predicted labels
    loss = K.abs(y_true - y_pred)
    loss = K.sum(loss, axis = 1)
    loss = K.sum(loss)
    return loss

def label_nn():
    # input layer
    img_input = Input(shape=(32, 32, 3))
    noisy_label = Input(shape = (10))

    # transfer learning - using VGG16 here
    base = VGG16(
        include_top=False,
        weights="imagenet",
        input_shape=(32,32,3),
        pooling='max'
    )

    # make the last layer of VGG16 trainable
    base.trainable = False
    base.get_layer('block5_conv3').trainable = True

    # use VGG16 as the base model
    img_vec = base(img_input)

    noisy_l = Dense(10)(noisy_label)
    img_vec = Dense(256)(img_vec)

    # concatenate noisy labels and image features
    x = Concatenate(axis=-1)([noisy_l, img_vec])
    x = Dense(256, activation = 'relu')(x)
    out = Dense(10, activation = 'softmax')(x)

    model = Model([img_input, noisy_label], out)
    # compile
    model.compile(loss=label_loss, metrics=['acc'], optimizer=RMSprop(0.
001))

    return model

```

```
In [46]: # record the computational time
start = time.time()
model = label_nn()
# train the label model
model.fit([clean_imgs, noisy], clean, batch_size = 128, epochs = 6)
end = time.time()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 3s 0us/step
58900480/58889256 [=====] - 3s 0us/step
Epoch 1/6
79/79 [=====] - 136s 2s/step - loss: 174.6149
- acc: 0.3415
Epoch 2/6
79/79 [=====] - 85s 1s/step - loss: 146.2020 -
acc: 0.4466
Epoch 3/6
79/79 [=====] - 82s 1s/step - loss: 130.1497 -
acc: 0.5086
Epoch 4/6
79/79 [=====] - 84s 1s/step - loss: 120.2542 -
acc: 0.5473
Epoch 5/6
79/79 [=====] - 100s 1s/step - loss: 114.5621
- acc: 0.5658
Epoch 6/6
79/79 [=====] - 94s 1s/step - loss: 110.4072 -
acc: 0.5818
```

```
In [47]: print(f"Total Label Network training Time: {end-start}")
```

```
Total Label Network training Time: 587.4192531108856
```

```
In [48]: # save model
model.save("../output/model_labelclean.h5")
```

```
In [49]: #Predict new labels for noisy set
noisy_imgs = imgs[10000:]/255
noisy_l = np.eye(10)[noisy_labels[10000:]]
```

```
In [50]: # record the computational time
start = time.time()
# predict labels
new_pred = model.predict([noisy_imgs, noisy_l])
end = time.time()
```

```
In [51]: print(f"Total Label Network prediction Time: {end-start}")
```

```
Total Label Network prediction Time: 443.62471890449524
```

```
In [58]: #clean up label vectors  
row_maxes = new_pred.argmax(axis=1)  
new_labels = np.eye(10)[row_maxes]  
  
#Create new train set from clean images and new pred labels  
upd_imgs = imgs  
upd_labels = np.concatenate((all_labels[:10000], new_labels , axis=0)
```

Train Model II with clean labels and new labels from label network

```
In [59]: # shuffle the images and split the data into training and validation set  
(0.9 and 0.1)  
shuff_imgs, target_vec = shuffle(upd_imgs, upd_labels, random_state=0)  
img_train, img_test, y_train, y_test = train_test_split(shuff_imgs, target_vec, test_size=0.10, random_state=42)
```

```
In [60]: # record the computational time
start = time.time()
# train model 2
modelIII = modelI()
modelIII, data_genII = train_model(modelIII, img_train, y_train, img_test,
y_test, "modelIII.h5", 10)
end = time.time()
```

```
Epoch 1/10
352/352 [=====] - 219s 618ms/step - loss: 1.67
63 - accuracy: 0.4041 - val_loss: 1.4858 - val_accuracy: 0.4860

Epoch 00001: val_accuracy improved from -inf to 0.48600, saving model to
o ../output/modelIII.h5
Epoch 2/10
352/352 [=====] - 153s 433ms/step - loss: 1.45
79 - accuracy: 0.4904 - val_loss: 1.4326 - val_accuracy: 0.4972

Epoch 00002: val_accuracy improved from 0.48600 to 0.49720, saving model
l to ../output/modelIII.h5
Epoch 3/10
352/352 [=====] - 135s 383ms/step - loss: 1.38
35 - accuracy: 0.5183 - val_loss: 1.3575 - val_accuracy: 0.5228

Epoch 00003: val_accuracy improved from 0.49720 to 0.52280, saving model
l to ../output/modelIII.h5
Epoch 4/10
352/352 [=====] - 139s 396ms/step - loss: 1.33
92 - accuracy: 0.5350 - val_loss: 1.3043 - val_accuracy: 0.5456

Epoch 00004: val_accuracy improved from 0.52280 to 0.54560, saving model
l to ../output/modelIII.h5
Epoch 5/10
352/352 [=====] - 128s 364ms/step - loss: 1.30
63 - accuracy: 0.5445 - val_loss: 1.2958 - val_accuracy: 0.5548

Epoch 00005: val_accuracy improved from 0.54560 to 0.55480, saving model
l to ../output/modelIII.h5
Epoch 6/10
352/352 [=====] - 164s 467ms/step - loss: 1.28
34 - accuracy: 0.5528 - val_loss: 1.2958 - val_accuracy: 0.5506

Epoch 00006: val_accuracy did not improve from 0.55480
Epoch 7/10
352/352 [=====] - 123s 348ms/step - loss: 1.27
31 - accuracy: 0.5566 - val_loss: 1.2914 - val_accuracy: 0.5464

Epoch 00007: val_accuracy did not improve from 0.55480
Epoch 8/10
352/352 [=====] - 138s 393ms/step - loss: 1.25
13 - accuracy: 0.5646 - val_loss: 1.2486 - val_accuracy: 0.5724

Epoch 00008: val_accuracy improved from 0.55480 to 0.57240, saving model
l to ../output/modelIII.h5
Epoch 9/10
352/352 [=====] - 121s 344ms/step - loss: 1.23
72 - accuracy: 0.5683 - val_loss: 1.2967 - val_accuracy: 0.5538

Epoch 00009: val_accuracy did not improve from 0.57240
Epoch 10/10
352/352 [=====] - 127s 360ms/step - loss: 1.22
83 - accuracy: 0.5699 - val_loss: 1.2420 - val_accuracy: 0.5626

Epoch 00010: val_accuracy did not improve from 0.57240
```

```
In [61]: print(f"Total Model II training Time: {end-start}")
```

Total Model II training Time: 1468.9489738941193

```
In [64]: # clean images and labels
```

```
img_cl = imgs[:10000]
```

```
y_cl = np.eye(10)[clean_labels]
```

```
# estimate the model accuracy using clean labels
```

```
cl_metrics = modelII.evaluate(data_genII.flow(img_cl, y_cl))
```

```
print(f"Model II Clean Labels Loss, Accuracy: {cl_metrics}")
```

313/313 [=====] - 9s 28ms/step - loss: 1.2218

- accuracy: 0.5703

Model II Clean Labels Loss, Accuracy: [1.2217928171157837, 0.5702999830245972]

```
In [8]: # [ADD WEAKLY SUPERVISED LEARNING FEATURE TO MODEL I]
```

```
def model_II(image):
```

```
    '''
```

*This function should takes in the image of dimension 32*32*3 as input and returns a label prediction*

```
    '''
```

```
    # load the model weights
```

```
    model2 = modelI()
```

```
    model2.load_weights("../output/modelII.h5")
```

```
    # predict
```

```
    pred = model2.predict(data_genII.flow(image))
```

```
    pred_class = np.argmax(pred, axis = 1)
```

```
    return pred_class
```

3. Evaluation

For assessment, we will evaluate your final model on a hidden test dataset with clean labels by the `evaluation` function defined as follows. Although you will not have the access to the test set, the function would be useful for the model developments. For example, you can split the small training set, using one portion for weakly supervised learning and the other for validation purpose.

```
In [9]: # [DO NOT MODIFY THIS CELL]
```

```
def evaluation(model, test_labels, test_imgs):
```

```
    y_true = test_labels
```

```
    y_pred = []
```

```
    for image in test_imgs:
```

```
        y_pred.append(model(image))
```

```
    print(classification_report(y_true, y_pred))
```

```
In [10]: # [DO NOT MODIFY THIS CELL]
# This is the code for evaluating the prediction performance on a testset
# You will get an error if running this cell, as you do not have the testset
# Nonetheless, you can create your own validation set to run the evaluation
n_test = 10000
test_labels = np.genfromtxt('../data/test_labels.csv', delimiter=',', dtype="int8")
test_imgs = np.empty((n_test, 32, 32, 3))
for i in range(n_test):
    img_fn = f'../data/test_images/test{i+1:05d}.png'
    test_imgs[i, :, :, :] = cv2.cvtColor(cv2.imread(img_fn), cv2.COLOR_BGR2RGB)
evaluation(baseline_model, test_labels, test_imgs)
```

	precision	recall	f1-score	support
0	0.33	0.46	0.38	1000
1	0.21	0.31	0.25	1000
2	0.20	0.04	0.07	1000
3	0.19	0.12	0.14	1000
4	0.24	0.48	0.32	1000
5	0.20	0.11	0.14	1000
6	0.24	0.34	0.28	1000
7	0.31	0.04	0.08	1000
8	0.27	0.43	0.33	1000
9	0.20	0.12	0.15	1000
accuracy			0.24	10000
macro avg	0.24	0.24	0.21	10000
weighted avg	0.24	0.24	0.21	10000

The overall accuracy is 0.24, which is better than random guess (which should have a accuracy around 0.10). For the project, you should try to improve the performance by the following strategies:

- Consider a better choice of model architectures, hyperparameters, or training scheme for the predictive model;
- Use both `clean_noisy_trainset` and `noisy_trainset` for model training via **weakly supervised learning** methods. One possible solution is to train a "label-correction" model using the former, correct the labels in the latter, and train the final predictive model using the corrected dataset.
- Apply techniques such as k -fold cross validation to avoid overfitting;
- Any other reasonable strategies.

4. Save Label Predictions CSV


```
In [29]: #Load trained Model 1/2 networks
model1 = modelI()
model1.load_weights("../output/modelI.h5")

model2 = modelI()
model2.load_weights("../output/modelII.h5")
```

```
In [30]: import glob
#Load test image data
test_fp = "../data/images/*"
test_img_files = glob.glob(test_fp)
test_img_files = test_img_files[:10000]

test_imgs = np.empty((10000,32,32,3))
i=0
for f in test_img_files:
    test_imgs[i,:,:,:]=cv2.cvtColor(cv2.imread(f),cv2.COLOR_BGR2RGB)
    i+=1
```

```
In [ ]: #Baseline Predictions
baseline_pred = []
for im in test_imgs:
    baseline_pred.append(baseline_model(im)[0])
```

```
In [ ]: #Model I
model1_pred = model1.predict(data_genI.flow(test_imgs))
model1_pred = np.argmax(model1_pred, axis = 1)
```

```
In [ ]: #Model II
model2_pred = model2.predict(data_genII.flow(test_imgs))
model2_pred = np.argmax(model2_pred, axis = 1)
```

```
In [ ]: #Save csv file
import pandas as pd

pred = read_csv("../output/label_prediction.csv")

pred['Baseline'] = baseline_pred
pred['Model I'] = model1_pred
pred["Model II"] = model2_pred

pred.to_csv("../output/label_prediction.csv")
```

Appendix

We ran our code on Google Colab and it works well with a RAM of 12.69 GB.

Here are the basic structures of some models we tried. We modified layers and different values of the parameters to see which structure has a higher accuracy.

```
In [ ]: # Artificial neural network (ANN)
# model=Sequential()
# model.add(Flatten(input_shape=(32,32,3)))
# model.add(Dense(256,activation='relu'))
# model.add(Dense(10,activation='softmax'))
```

```
In [ ]: # A multilayer perceptron (MLP)
# model = Sequential()
# model.add(Dense(256, activation='relu', input_dim=3072))
# model.add(Dense(256, activation='relu'))
# model.add(Dense(10, activation='softmax'))
```

```
In [ ]: # Convolutionary neural network(CNN)
# model = Sequential()
# model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32,
3)))
# model.add(MaxPooling2D(pool_size=(2, 2)))
# model.add(Flatten())
# model.add(Dense(256, activation='relu'))
# model.add(Dense(10, activation='softmax'))
```

```

In [ ]: # Transfer Learning: VGG 16, VGG 19, ResNet50, GoogLeNet and ArcFace
# for example VGG19
# base_model = VGG19(include_top=False,weights='imagenet',input_shape=(3
2,32,3),classes=y_train.shape[1])
# base_model.trainable = False
# img_input = Input(shape=(32, 32, 3))
# model = base_model(img_input)
# model = Flatten()(model)
# model = Dense(512,activation=('relu'))(model)
# out = Dense(10,activation=('softmax'))(model)
# model = Model(img_input, out)

# for example ResNet50
# base_model = ResNet50(include_top=False,weights='imagenet',input_shape
=(224,224,3),pooling='max')
# base_model.trainable = False
# img_input = Input(shape=(32,32,3))
# model = UpSampling2D(size=(7,7))(img_input)
# model = base_model_2(model)
# model = Flatten()(model)
# model = Dense(512, activation="relu")(model)
# out = Dense(10, activation="softmax")(model)
# model = Model(img_input, out)

# for example ArcFace

# base_model = ArcFaceModel(size=32, channels=3, num_classes=None, name
='arcface_model',
                        #margin=0.5, logist_scale=64, embd_shape=512,
                        #head_type='ArcHead', backbone_type='ResNet50',
                        #w_decay=5e-4, use_pretrain=True, training=False)
# base_model.trainable = False
# img_input = Input(shape=(32, 32, 3))
# model = base_model(img_input)
# model = Flatten()(model)
# model = BatchNormalization()(model)
# model = Dense(256, activation='relu')(model)
# model = Dropout(0.3)(model)
# model = BatchNormalization()(model)
# model = Dense(128, activation='relu')(model)
# model = Dropout(0.3)(model)
# model = BatchNormalization()(model)
# model = Dense(64, activation='relu')(model)
# model = Dropout(0.3)(model)
# out = Dense(10, activation='softmax')(model)
# model = Model(img_input, out)

```

To reduce overfitting, we tried some layers and methods

```
In [ ]: # for layers

# Modify the parameter "activity_regularizer" of the dense layer
# activity_regularizer = regularizers.l2(0.01)

# Dropout layer randomly sets input units to 0 at the given rate for each step during training
# Dropout(0.25)

# Batch Normalization layer normalizes the inputs
# BatchNormalization()

# for methods
# Early stopping stops training when a metric stops improving
# early = EarlyStopping(monitor='loss', patience=3)

# when we are using transfer learning the overfitting is about 8%
# we make some layers of the base model (such as the last fully connected layer)
# trainable in order to reduce overfitting
# base_model.get_layer('block5_conv4').trainable = True
```

We tried different optimizers and adjust the value of parameters

```
In [ ]: # Gradient descent (with momentum) optimizer
# sgd = SGD(learning_rate=0.001, momentum=.9, nesterov=False)

# RMSprop
# rms = RMSprop(learning_rate=0.001, rho=0.9, momentum=0.0, epsilon=1e-07, centered=False)

# Adam
# adam = Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, amsgrad=False)

# Nadam
# nadam = Nadam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07)

# Adadelta
# ada = Adadelta(learning_rate=0.1, rho=0.95, epsilon=1e-07)
```

For the learning rate of the optimizer, we tried learning rate schedulers

```
In [ ]: # ExponentialDecay
# an exponential decay schedule
# lr_schedule = ExponentialDecay(initial_learning_rate=1e-2, decay_steps
=10000, decay_rate=0.90)

# ReduceLRonPlateau
# this reduces learning rate when a metric stops improving
# lrr= ReduceLRonPlateau( monitor='val_accuracy', factor=.01, patience=
3, min_lr=1e-5)
# callbacks = [lrr]
```

After adjusting model structure, tuning the values of hyperparameters and applying different methods, model I and model II are the best of all. The accuracy of ANN, MLP and CNN with other structures is about low 20s. The accuracy of transfer learning is about 35s but takes almost ten minutes per epoch.